# Transactions

COSC 304 – Introduction to Database Systems

# Transaction Management Overview

The database system must ensure that the data is always *consistent*. Two challenges in preserving consistency:

- 1) Must handle *failures* of various kinds (hardware, crashes).
- 2) Must support *concurrent execution* of multiple transactions and guarantee that concurrency does not cause inconsistency.

A *transaction* is an *atomic* program that executes on the database and preserves the consistency of the database.

- The input to a transaction is a consistent database AND the output of the transaction must also be a consistent database.
- A transaction must execute completely or not at all.

# Transaction Management - Motivating Example

Consider a person who wants to transfer $50 from a savings account with balance $1000 to a checking account with current balance $250.

- 1) At the ATM, the person starts the process by telling the bank to remove $50 from the savings account.
- 2) The $50 is removed from the savings account by the bank.
- 3) Before the customer can tell the ATM to deposit the $50 in the checking account, the ATM "crashes."

Where has the $50 gone?

It is lost if the ATM did not support transactions!
The customer wanted the withdraw and deposit to both happen in one step, or neither action to happen.

# ACID Properties

To preserve integrity, transactions have the following properties:

- *Atomicity* -  Either all operations of the transaction are properly reflected in the database or none are.
- *Consistency* -  Execution of a transaction in isolation preserves the consistency of the database.
- *Isolation* -  Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions.
- *Durability* -  After a transaction successfully completes, the changes it has made to the database persist, even if there are system failures.

# ACID Properties

*Question:* Two transactions running at the same time can see each other's updates.  What ACID property is violated?

**A)** atomicity

**B)** consistency

**C)** isolation

**D)** durability

**E)** none of them

# Transaction Definition in SQL

In SQL, a transaction begins implicitly.

A transaction in SQL ends by:
- **Commit** accepts updates of current transaction.
- **Rollback** aborts current transaction and discards its updates. Failures may also cause a transaction to be aborted.

In programming with databases you can often set the connection to **auto-commit** which means a commit is done after every statement.

# Example Transactions

Transaction to deposit $50 into a bank account:

```
-- TRANSACTION T1:
UPDATE Account WHERE num = 'S1' SET balance=balance+50;
COMMIT;
```

Transaction to calculate totals for all accounts (twice):

```
-- TRANSACTION T2:
SELECT SUM(balance) as total1 FROM Account;
SELECT SUM(balance) as total2 FROM Account;
COMMIT;
```

Transaction to add a new account:

```
-- TRANSACTION T3:
INSERT INTO ACCOUNT (num, balance) VALUES ('S5', 100);
COMMIT;
```

# Isolation Levels in SQL

An ***isolation level*** reflects how a transaction perceives the results of other transactions. Lowering isolation level may improve performance but may sacrifice consistency. The isolation level can be specified by:

`SET TRANSACTION ISOLATION LEVEL = X`  where $X$ is

- **Serializable -** transactions behave like executed one at a time.
- **Repeatable read -** repeated reads must return same data.  Does not necessarily read newly inserted records.   Affected by insert
- **Read committed -** only committed values can be read, but successive reads may return different values.
- **Read uncommitted -** even uncommitted records may be read.  Reading an uncommitted value is called a *dirty read*.

# Scheduling of Transactions

Each transaction in a database is a separate executing program.

- A transaction may be its own program or a thread of execution.

The operating system schedules the execution of programs outside of the control of the DBMS.

- Transactions may be executed in any order (as long as the order of operations within a transaction are the same).  This interleaving produces different schedules.

The DBMS uses its concurrency control protocol to restrict the schedules to those that respect the consistency specified by the user for the transaction isolation level.

- All transactions must write lock any data item updated and lock the relation if inserting.
- Isolation level only affects read locks.

# Transaction Schedule Example

Transaction **T2** that does two queries has its statements interleaved with transaction **T1** that does an update:

```
SELECT SUM(balance) as total1 FROM Account;  --T2
UPDATE Account WHERE num='S1' SET balance=balance+50;--T1
COMMIT;  -- T1
SELECT SUM(balance) as total2 FROM Account; -- T2
COMMIT; -- T2
```

With isolation level *read committed*, `total1` will not be the same as `total2`.

With isolation level *repeatable read*, this schedule is *not possible* as **T2** will lock all accounts stopping **T1** from updating.

Transaction **T2** that does two queries has its statements interleaved with transaction **T3** that does an insert:

```
SELECT SUM(balance) as total1 FROM Account;  --T2
INSERT INTO ACCOUNT (num, balance) VALUES ('S5' , 100);
COMMIT;  -- T3
SELECT SUM(balance) as total2 FROM Account; -- T2
COMMIT; -- T2
```

With isolation level *repeatable read*, `total1` will not be the same as `total2`.

With isolation level *serializable*, this schedule is not possible as **T2** will lock the account table, stopping **T3** from inserting.

# Summary of Isolation Levels

| Isolation Level | Problems | Lock Usage | Speed | Comments |
|---|---|---|---|---|
| Serializable | None | Read locks held to commit ; read lock on relation | Slowest | Only level that guarantees correctness. |
| Repeatable read | Phantom tuples | Read locks held to commit | Medium | Useful for modify transactions. May not see tuples inserted by others. |
| Read committed | Phantom tuples, values may change | Read locks released after each statement | Fast | Useful for transactions where operations are separable but updates are all or none. Re-reading same value may produce different results. |
| Read uncommitted | Phantoms, values may change, dirty reads | No read locks | Fastest | Useful for read-only transactions that tolerate inaccurate results. May see updates that will never be committed. |

# Read Committed Question

*Question:* Will this transaction always see the same results for both queries if executed using **READ COMMITTED**?

```
SELECT SUM(balance) as total1 FROM Account;
SELECT SUM(balance) as total2 FROM Account;
COMMIT;
```

**A)** yes

**B)** no

*Question:* Will this transaction always see the same results for both queries if executed using **REPEATABLE READ**?

```
SELECT SUM(balance) as total1 FROM Account;
SELECT SUM(balance) as total2 FROM Account;
COMMIT;
```

**A)** yes

**B)** no

# JDBC Transaction Example

```
try (Connection con = DriverManager.getConnection(url, uid, pw);
        Statement stmt = con.createStatement();)
{
    con.setAutocommit(false);        ⟵ Force explicit commit/rollback
    ResultSet rst = stmt.executeQuery("SELECT ename,salary
                                        FROM emp WHERE eno='E1'");
    if (rst.next() && rst.getDouble(2) < 50000)
    {
        stmt.executeUpdate("UPDATE emp SET salary=100000 WHERE eno='E1'");
        con.commit();                ⟵ Commit work to DB
    } else
        con.rollback();              ⟵ Rollback work done
}
catch (SQLException ex) {
    System.err.println(ex); con.rollback();
}
```

# Recursive Queries in SQL

General form:

```
WITH RECURSIVE tableName(attr1,attr2,..attrN) AS
    <SELECT query that defines recursive relation>
<SELECT query that uses recursive relation>
```

Example: Return all employees supervised by 'J. Jones'.

```
WITH RECURSIVE supervises(supId,empId) AS (
    SELECT supereno, eno FROM emp WHERE supereno IS NOT NULL
    UNION ALL
    SELECT S.supId, E.eno
    FROM supervises S, emp E
    WHERE S.empId = E.supereno )
SELECT E.eno, E.ename FROM supervises, emp AS S, emp E
WHERE S.ename = 'J. Jones' and supervises.supId = S.eno
        and supervises.empId = E.eno;
```

# Conclusion

A *transaction* is an *atomic* program that executes on the database and preserves the consistency of the database.

In SQL, different isolation levels can be specified:

- serializable, repeatable read, read committed, read uncommitted
- Weaker forms of isolation do not guarantee the ACID properties, but may be useful for read transactions that require faster execution.

Object-relational DBMSs support user defined data types, active consistency checks (triggers), inheritance, and recursive queries.

- `WITH RECURSIVE` syntax is used to write recursive SQL queries.

# Objectives

- Define: transaction
- List ACID properties.
- List and explain the isolation levels in SQL
- Be aware of WITH RECURSIVE for recursive SQL queries.